

## ▼ 1a (i)

```
from keras.datasets import mnist
import tensorflow as tf
import numpy as np
import matplotlib.pyplot as plt
import random
import pandas as pd
import seaborn as sns
```

```
(x_train, y_train), (x_test, y_test) = mnist.load_data()
```

```
x_train.shape
```

```
(60000, 28, 28)
```

```
y_train.shape
```

```
(60000,)
```

```
y_train
```

```
array([5, 0, 4, ..., 5, 6, 8], dtype=uint8)
```

```
from google.colab.patches import cv2_imshow
cv2_imshow(x_train[0])
cv2_imshow(x_train[2])
```



```
my_list = np.array(y_train)
indices_all=[]
for i in range(0,10):
    indices = np.where(my_list == i)[0]
    indices = indices.tolist()
    indices_all.append(indices)
# display result
print(indices_all)
```

```
[[1, 21, 34, 37, 51, 56, 63, 68, 69, 75, 81, 88, 95, 108, 114, 118, 119, 121,
```



```
n = 5
indices_store=[]
for i in range(0,10) :
    indices_store.append(random.sample(indices_all[i], n))
```

indices\_store

```
[[30187, 51669, 18093, 53290, 53424],
 [38538, 41282, 21357, 49680, 16914],
 [27451, 18206, 44226, 48877, 25150],
 [12920, 31140, 41048, 51701, 54285],
 [33554, 50656, 31111, 14027, 41020],
 [31584, 7848, 44713, 56695, 16558],
 [25171, 37829, 5141, 3736, 46632],
 [21693, 23127, 16809, 20105, 22350],
 [53737, 47498, 19582, 34933, 36974],
 [32848, 8510, 27440, 19914, 49227]]
```

indices\_store[0][0]

30187

```
for i in range(0,10):
    for c in range(0,5):
        cv2_imshow(x_train[indices_store[i][c]])
```





1a(ii)(A)

```

x_train.shape

```

```

(60000, 28, 28)

```

```

x_test.shape

```

```

(10000, 28, 28)

```

```


```

```


```

We observe in trainX, there are 60000 images of shape 28x28 and in testX there are 10000 images of shape 28x28. Thus, the shape of all images is 28x28. Hence, yes, all images are of the same size.

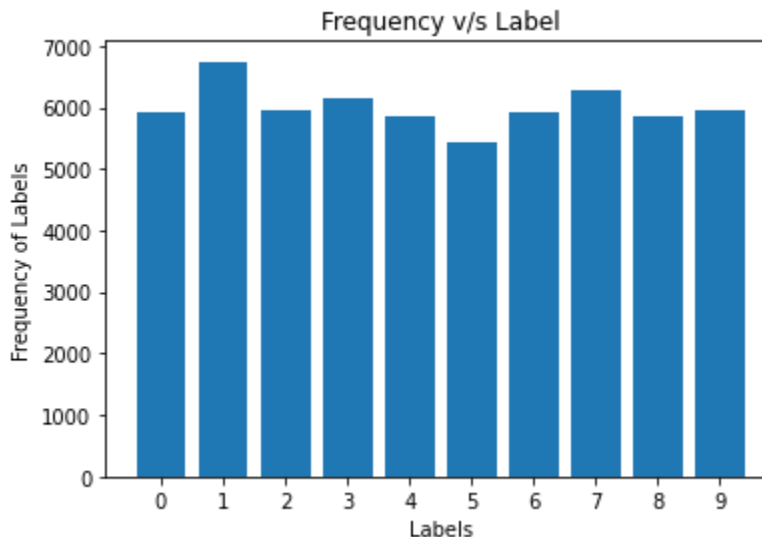


## 1a(ii)(B)

```
(x_train, y_train), (x_test, y_test) = mnist.load_data()

labels, count = np.unique(y_train, return_counts=True)

plt.bar(labels, count, align='center')
plt.xticks(labels)
plt.xlabel('Labels')
plt.ylabel('Frequency of Labels')
plt.title('Frequency v/s Label')
plt.show()
```



Since there is no significant difference in the frequencies, there is no substantial class imbalance as the number of images for all labels is approximately equal to 6000.

## 1a(ii)(C)

```
(x_train, y_train), (x_test, y_test) = mnist.load_data()
print("x_train before normalization: ", x_train[0])
x_train = x_train.astype('float32')/255
x_test = x_test.astype('float32')/255
print("x_train after normalization: ", x_train[0])
```

[illegible]

```

[[0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]]
x_train after normalization: [[0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]]

```

We can normalise `x_train` by dividing it by the maximum pixel value, i.e. 255. To visualise, I have printed the first element of `x_train` before and after normalisation. The value before normalisation ranges from 0 to 255, whereas after normalisation it ranges from 0 to 1.

## 1b

```

from google.colab import drive
drive.mount('/content/drive')

```

Drive already mounted at `/content/drive`; to attempt to forcibly remount, call

```

data = pd.read_csv("/content/drive/MyDrive/Datasets/city_day.csv")
data.head()

```

	City	Date	PM2.5	PM10	NO	NO2	NOx	NH3	CO	S02	O3
0	Ahmedabad	2015-01-01	NaN	NaN	0.92	18.22	17.15	NaN	0.92	27.64	133.36
1	Ahmedabad	2015-01-02	NaN	NaN	0.97	15.69	16.46	NaN	0.97	24.55	34.06
2	Ahmedabad	2015-01-03	NaN	NaN	17.40	19.30	29.70	NaN	17.40	29.07	30.70
3	Ahmedabad	2015-01-04	NaN	NaN	1.70	18.48	17.97	NaN	1.70	18.59	36.08
4	Ahmedabad	2015-01-05	NaN	NaN	22.10	21.42	37.76	NaN	22.10	39.33	39.31

```


data.hist(column=["NO"])
data.hist(column=["NO2"])
data.hist(column=["NOx"])
data.hist(column=["CO"])
data.hist(column=["S02"])
data.hist(column=["O3"])
data.hist(column=["Benzene"])
data.hist(column=["Toluene"])
data.hist(column=["Xylene"])


```

```


array([[<matplotlib.axes._subplots.AxesSubplot object at 0x7f851a15f640>]],
      dtype=object)

```







NO2



NOx



CO



```
data['Year'] = pd.to_datetime(data['Date']).dt.year
data_new = data[(data.Year >= 2015) & (data.Year <= 2018)]
print(data_new.describe())
```

	PM2.5	PM10	NO	N02	N0x	\
count	13319.000000	7961.000000	14300.000000	14354.000000	14029.000000	
mean	79.914880	137.369947	18.282866	31.070509	33.021959	
std	74.004494	100.585837	23.589004	24.744464	33.528014	
min	1.720000	0.010000	0.020000	0.010000	0.000000	
25%	34.695000	69.710000	5.960000	13.900000	11.830000	
50%	57.380000	107.770000	10.160000	24.575000	23.330000	
75%	96.200000	170.790000	20.040000	40.890000	41.800000	
max	949.990000	917.080000	287.140000	362.210000	467.630000	

	NH3	CO	S02	O3	Benzene	\
count	9141.000000	15905.000000	14100.000000	13995.000000	13619.000000	
mean	28.458301	2.530362	14.435597	35.187022	2.585403	
std	31.722477	7.686014	19.112619	23.047512	8.866921	
min	0.010000	0.000000	0.040000	0.010000	0.000000	
25%	10.390000	0.430000	5.170000	18.515000	0.090000	
50%	20.090000	0.940000	8.390000	30.690000	0.850000	
75%	35.480000	1.600000	14.842500	46.360000	2.940000	
max	352.890000	175.810000	193.860000	257.730000	391.880000	

	Toluene	Xylene	AQI	Year
count	13103.000000	6803.000000	13358.000000	17439.000000
mean	7.461200	2.990104	189.250262	2016.850393
std	14.545896	5.934221	152.539902	1.091172
min	0.000000	0.000000	13.000000	2015.000000
25%	0.435000	0.020000	91.000000	2016.000000
50%	2.570000	0.650000	136.000000	2017.000000
75%	7.950000	3.230000	253.000000	2018.000000
max	411.520000	125.180000	2049.000000	2018.000000

0      50      100      150      200      250

```
print("Number of null values in the dataset: ",(data_new.isnull().sum().sum()))
```

Number of null values in the dataset: 66801

~~~~~ |  | | | | |

```
#Normalization
```

```
for i in data.columns:
```

```
    if i not in ["City", "Date", "AQI_Bucket"]:
```

```
        data[i] = (data[i] - data[i].min()) / (data[i].max())
```

```
data
```

|   | City      | Date       | PM2.5 | PM10 | NO       | N02      | N0x      | N   |
|---|-----------|------------|-------|------|----------|----------|----------|-----|
| 0 | Ahmedabad | 2015-01-01 | NaN   | NaN  | 0.002304 | 0.050275 | 0.036674 | NaN |
| 1 | Ahmedabad | 2015-01-02 | NaN   | NaN  | 0.002432 | 0.043290 | 0.035199 | NaN |
| 2 | Ahmedabad | 2015-01-03 | NaN   | NaN  | 0.044487 | 0.053256 | 0.063512 | NaN |
| 3 | Ahmedabad | 2015-01-04 | NaN   | NaN  | 0.004300 | 0.050993 | 0.038428 | NaN |
| 4 | Ahmedabad | 2015-01-05 | NaN   | NaN  | 0.056517 | 0.059109 | 0.080748 | NaN |



```

...
...
...
...
...
...
...
...
29526  Visakhapatnam  2020-06-27  0.015769  0.05093  0.019607  0.069159  0.041785  0.0353
29527  Visakhapatnam  2020-06-28  0.025621  0.07408  0.008703  0.071920  0.035348  0.0339
29528  Visakhapatnam  2020-06-29  0.024074  0.06572  0.008780  0.081500  0.039198  0.0303
29529  Visakhapatnam  2020-06-30  0.017474  0.04996  0.010315  0.080754  0.040203  0.0283
29530  Visakhapatnam  2020-07-01  0.015748  0.06599  0.000973  0.074101  0.030045  0.0147

```

29531 rows x 17 columns

```
| | | | | | | |
```

data.info()

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 29531 entries, 0 to 29530
Data columns (total 17 columns):
#   Column                Non-Null Count  Dtype
---  -
0   City                   29531 non-null  object
1   Date                   29531 non-null  object
2   PM2.5                  24933 non-null  float64
3   PM10                   18391 non-null  float64
4   NO                      25949 non-null  float64
5   NO2                    25946 non-null  float64
6   NOx                     25346 non-null  float64
7   NH3                    19203 non-null  float64
8   CO                      27472 non-null  float64
9   SO2                     25677 non-null  float64
10  O3                      25509 non-null  float64
11  Benzene                 23908 non-null  float64
12  Toluene                 21490 non-null  float64
13  Xylene                  11422 non-null  float64
14  AQI                     24850 non-null  float64
15  AQI_Bucket              24850 non-null  object
16  Year                    29531 non-null  float64
dtypes: float64(14), object(3)
memory usage: 3.8+ MB

```

### 3c: Linear Regression

```

import pandas as pd
from sklearn import datasets
from sklearn import linear_model
from sklearn import model_selection
from sklearn.model_selection import train_test_split
# Load the dataset
from sklearn.model_selection import KFold
diabetes = datasets.load_diabetes()

```

```
diabetes = datasets.load_diabetes()
from matplotlib import pyplot as plt
import numpy as np
```

```
diabetes = datasets.load_diabetes(as_frame=True)
```

```
print(type(diabetes['data']))
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
data = pd.read_csv("/content/drive/MyDrive/Datasets/diabetes-dataset.csv")
```

```
data
```

|      | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI  | Diabete |
|------|-------------|---------|---------------|---------------|---------|------|---------|
| 0    | 2           | 138     | 62            | 35            | 0       | 33.6 |         |
| 1    | 0           | 84      | 82            | 31            | 125     | 38.2 |         |
| 2    | 0           | 145     | 0             | 0             | 0       | 44.2 |         |
| 3    | 0           | 135     | 68            | 42            | 250     | 42.3 |         |
| 4    | 1           | 139     | 62            | 41            | 480     | 40.7 |         |
| ...  | ...         | ...     | ...           | ...           | ...     | ...  |         |
| 1995 | 2           | 75      | 64            | 24            | 55      | 29.7 |         |
| 1996 | 8           | 179     | 72            | 42            | 130     | 32.7 |         |
| 1997 | 6           | 85      | 78            | 0             | 0       | 31.2 |         |
| 1998 | 0           | 129     | 110           | 46            | 130     | 67.1 |         |
| 1999 | 2           | 81      | 72            | 15            | 76      | 30.1 |         |

2000 rows x 9 columns

```
df_max_scaled = data.copy()
```

```
# apply normalization techniques
```

```
for column in df_max_scaled.columns:
```

```
    df_max_scaled[column] = df_max_scaled[column] / df_max_scaled[column].abs().max()
```

```
df_max_scaled
```

|   | Pregnancies | Glucose  | BloodPressure | SkinThickness | Insulin  | BMI      | Dial |
|---|-------------|----------|---------------|---------------|----------|----------|------|
| 0 | 0.117647    | 0.693467 | 0.508197      | 0.318182      | 0.000000 | 0.416873 |      |
| 1 | 0.000000    | 0.422111 | 0.672131      | 0.281818      | 0.168011 | 0.473945 |      |

|      |          |          |          |          |          |          |
|------|----------|----------|----------|----------|----------|----------|
| 2    | 0.000000 | 0.728643 | 0.000000 | 0.000000 | 0.000000 | 0.548387 |
| 3    | 0.000000 | 0.678392 | 0.557377 | 0.381818 | 0.336022 | 0.524814 |
| 4    | 0.058824 | 0.698492 | 0.508197 | 0.372727 | 0.645161 | 0.504963 |
| ...  | ...      | ...      | ...      | ...      | ...      | ...      |
| 1995 | 0.117647 | 0.376884 | 0.524590 | 0.218182 | 0.073925 | 0.368486 |
| 1996 | 0.470588 | 0.899497 | 0.590164 | 0.381818 | 0.174731 | 0.405707 |
| 1997 | 0.352941 | 0.427136 | 0.639344 | 0.000000 | 0.000000 | 0.387097 |
| 1998 | 0.000000 | 0.648241 | 0.901639 | 0.418182 | 0.174731 | 0.832506 |
| 1999 | 0.117647 | 0.407035 | 0.590164 | 0.136364 | 0.102151 | 0.373449 |

2000 rows x 9 columns

```

data = df_max_scaled
train, test = train_test_split(data, test_size=0.1)

import math
class Linear_Regression:
    # Initiliasie the class with user defined learning rate and number of epochs
    def __init__(self, learning_rate, n_epochs):
        self.learning_rate = learning_rate
        self.n_epochs = n_epochs

    def gradient_descent(self,x,y):
        theta_curr = np.zeros(x.shape[1]) #weights
        bias = 0 #bias
        n_samples = y.size
        cost_l = [0] * self.n_epochs
        rmse = [0] * self.n_epochs

        for i in range(self.n_epochs):
            y_pred = x.dot(theta_curr) + bias

            #computing gradients
            dm = (-2/n_samples)*x.T.dot(y-y_pred)
            dc = (-2/n_samples)*np.sum(y-y_pred)

            #Updating weights
            theta_curr = theta_curr - self.learning_rate*dm
            bias = bias - self.learning_rate*dc

            #Calculating the cost associated with each iteration
            cost = (1/(2*n_samples))*np.sum((y_pred-y)**2)
            cost_l[i] = cost

```

```

        cost_curr = cost
        mse = ((1/(n_samples))*np.sum((y_pred-y)**2))
        mse = math.sqrt(mse)
        rmse[i] = mse
    return theta_curr, bias, cost_l, rmse

def predict(self,X,theta,bias):
    return X.dot(theta) + bias

model = LinearRegression(learning_rate=0.01,n_epochs=200)
Y= train['Outcome']
X= train.drop(['Outcome'],axis=1)

X=X.to_numpy()
Y=Y.to_numpy()


kf = KFold(n_splits=3)


index =0
rmse_sum =[]
t_min=[]

for train_index, test_index in kf.split(train):
    X_train, X_test = X[train_index], X[test_index]
    Y_train, Y_test = Y[train_index], Y[test_index]

    #t: Theta obtained after gradient descent, b: Bias obtained after gradient descent
    t,b,c,rmse = model.gradient_descent(X_train,Y_train)
    y_pred_train = X_train.dot(t)+b
    y_pred = model.predict(X_test,t,b) #validation set
    # fig, axes = plt.subplots(1, 3)
    # axes[0].plot(rmse,'g')
    plt.plot(rmse)
    t_len = len(t)
    t_min.extend(t)
    rmse_sum.append(sum(rmse))
min_ind = rmse_sum.index(min(rmse_sum))
t_best = t_min[min_ind*t_len:(min_ind*t_len+t_len)]

```





Performed 3 fold k cross validation to obtain better results


```
Y_test = test['Outcome']
X_test = test.drop(['Outcome'],axis=1)
Y_test = Y_test.to_numpy()
X_test = X_test.to_numpy()

y_pred_test = X_test.dot(t_best)

t,b,c,rmse = model.gradient_descent(X_test,Y_test)
y_pred_train = X_test.dot(t_best)+b
y_pred = model.predict(X_test,t_best,b)
plt.plot(rmse)

print("RMSE using Linear Regression from Scratch: ",min(rmse))
```

RMSE using Linear Regression from Scratch: 0.44993638716413276



```
from sklearn.linear_model import LinearRegression
from sklearn import metrics
regr = LinearRegression()
rmse_sci= []
t_reg = []
kfold_datasets_stored=[]

for train_index, test_index in kf.split(train):
    X_train, X_test = X[train_index], X[test_index]
    Y_train, Y_test = Y[train_index], Y[test_index]
    kfold_datasets_stored.append([X_train, X_test,Y_train, Y_test])
    regr = LinearRegression()
    regr.fit(X_train, Y_train)
    t_reg.extend(regr.coef )
```

```

    regr.extend(regr_list_)
    y_pred = regr.predict(X_test)
    rmse_sci.append(np.sqrt(metrics.mean_squared_error(Y_test,y_pred)))

min_index = np.argmin(np.array(rmse_sci))
X_train_kf, X_test_kf, Y_train_kf, Y_test_kf= kfold_datasets_stored[min_index]
model=LinearRegression()
model.fit(X_train_kf,Y_train_kf)
y_pred_kf= model.predict(X_test_kf)
np.sqrt(metrics.mean_squared_error(Y_test_kf,y_pred_kf))
rmse_sci = np.sqrt(metrics.mean_squared_error(Y_test_kf,y_pred_kf))
print("RMSE using Linear Regression (SciKit Learn): ", rmse_sci)

```

RMSE using Linear Regression (SciKit Learn): 0.39157319306864125

RMSE using Linear Regression (SciKit Learn) is 0.3974252185078301

RMSE using Linear Regression from Scratch is 0.4552361135405836

## Q4: Naive Bayes from Scratch

```

import numpy as np
from sklearn import datasets
from sklearn.model_selection import train_test_split

# Load Iris dataset
iris = datasets.load_iris()
data1 = pd.DataFrame(data= np.c_[iris['data'], iris['target']],
                    columns= iris['feature_names'] + ['target'])

data1

```

|     | sepal length<br>(cm) | sepal width<br>(cm) | petal length<br>(cm) | petal width<br>(cm) | target |
|-----|----------------------|---------------------|----------------------|---------------------|--------|
| 0   | 5.1                  | 3.5                 | 1.4                  | 0.2                 | 0.0    |
| 1   | 4.9                  | 3.0                 | 1.4                  | 0.2                 | 0.0    |
| 2   | 4.7                  | 3.2                 | 1.3                  | 0.2                 | 0.0    |
| 3   | 4.6                  | 3.1                 | 1.5                  | 0.2                 | 0.0    |
| 4   | 5.0                  | 3.6                 | 1.4                  | 0.2                 | 0.0    |
| ... | ...                  | ...                 | ...                  | ...                 | ...    |
| 145 | 6.7                  | 3.0                 | 5.2                  | 2.3                 | 2.0    |
| 146 | 6.3                  | 2.5                 | 5.0                  | 1.9                 | 2.0    |
| 147 | 6.5                  | 2.0                 | 5.2                  | 2.0                 | 2.0    |

|     |     |     |     |     |     |
|-----|-----|-----|-----|-----|-----|
| 147 | 6.5 | 3.0 | 5.2 | 2.0 | 2.0 |
| 148 | 6.2 | 3.4 | 5.4 | 2.3 | 2.0 |
| 149 | 5.9 | 3.0 | 5.1 | 1.8 | 2.0 |

150 rows x 5 columns

```
X = data1.iloc[:, :4]
y = data1.iloc[:, 4:]
print(X.shape, y.shape)
print(X)
```

```
(150, 4) (150, 1)
      sepal length (cm)  sepal width (cm)  petal length (cm)  petal width (cm)
0                5.1           3.5           1.4           0.2
1                4.9           3.0           1.4           0.2
2                4.7           3.2           1.3           0.2
3                4.6           3.1           1.5           0.2
4                5.0           3.6           1.4           0.2
..                ...           ...           ...           ...
145              6.7           3.0           5.2           2.3
146              6.3           2.5           5.0           1.9
147              6.5           3.0           5.2           2.0
148              6.2           3.4           5.4           2.3
149              5.9           3.0           5.1           1.8
```

[150 rows x 4 columns]

```
data1.describe()
```

|              | sepal length<br>(cm) | sepal width<br>(cm) | petal length<br>(cm) | petal width<br>(cm) | target     |
|--------------|----------------------|---------------------|----------------------|---------------------|------------|
| <b>count</b> | 150.000000           | 150.000000          | 150.000000           | 150.000000          | 150.000000 |
| <b>mean</b>  | 5.843333             | 3.057333            | 3.758000             | 1.199333            | 1.000000   |
| <b>std</b>   | 0.828066             | 0.435866            | 1.765298             | 0.762238            | 0.819232   |
| <b>min</b>   | 4.300000             | 2.000000            | 1.000000             | 0.100000            | 0.000000   |
| <b>25%</b>   | 5.100000             | 2.800000            | 1.600000             | 0.300000            | 0.000000   |
| <b>50%</b>   | 5.800000             | 3.000000            | 4.350000             | 1.300000            | 1.000000   |
| <b>75%</b>   | 6.400000             | 3.300000            | 5.100000             | 1.800000            | 2.000000   |
| <b>max</b>   | 7.900000             | 4.400000            | 6.900000             | 2.500000            | 2.000000   |

```
data1.dtypes
```

```
sepal length (cm)    float64
sepal width (cm)     float64
petal length (cm)    float64
petal width (cm)     float64
```

```

petal width (cm)      float64
target                float64
dtype: object

```

```
data1_notarget = data1.drop(['target'],axis=1)
```

```
data1.count()
```

```

sepal length (cm)      150
sepal width (cm)       150
petal length (cm)      150
petal width (cm)       150
target                 150
dtype: int64

```

```
data1[(data1.max() - data1.min()).idxmax()]
```

```

0      1.4
1      1.4
2      1.3
3      1.5
4      1.4
...
145    5.2
146    5.0
147    5.2
148    5.4
149    5.1
Name: petal length (cm), Length: 150, dtype: float64

```

```

cols = ['sepal length (cm)', 'sepal width (cm)', 'petal length (cm)', 'petal width
for col in cols:
    print(data1[col].unique())

```

```

[5.1 4.9 4.7 4.6 5.  5.4 4.4 4.8 4.3 5.8 5.7 5.2 5.5 4.5 5.3 7.  6.4 6.9
 6.5 6.3 6.6 5.9 6.  6.1 5.6 6.7 6.2 6.8 7.1 7.6 7.3 7.2 7.7 7.4 7.9]
[3.5 3.  3.2 3.1 3.6 3.9 3.4 2.9 3.7 4.  4.4 3.8 3.3 4.1 4.2 2.3 2.8 2.4
 2.7 2.  2.2 2.5 2.6]
[1.4 1.3 1.5 1.7 1.6 1.1 1.2 1.  1.9 4.7 4.5 4.9 4.  4.6 3.3 3.9 3.5 4.2
 3.6 4.4 4.1 4.8 4.3 5.  3.8 3.7 5.1 3.  6.  5.9 5.6 5.8 6.6 6.3 6.1 5.3
 5.5 6.7 6.9 5.7 6.4 5.4 5.2]
[0.2 0.4 0.3 0.1 0.5 0.6 1.4 1.5 1.3 1.6 1.  1.1 1.8 1.2 1.7 2.5 1.9 2.1
 2.2 2.  2.4 2.3]

```

```
iris = sns.load_dataset("iris")
```



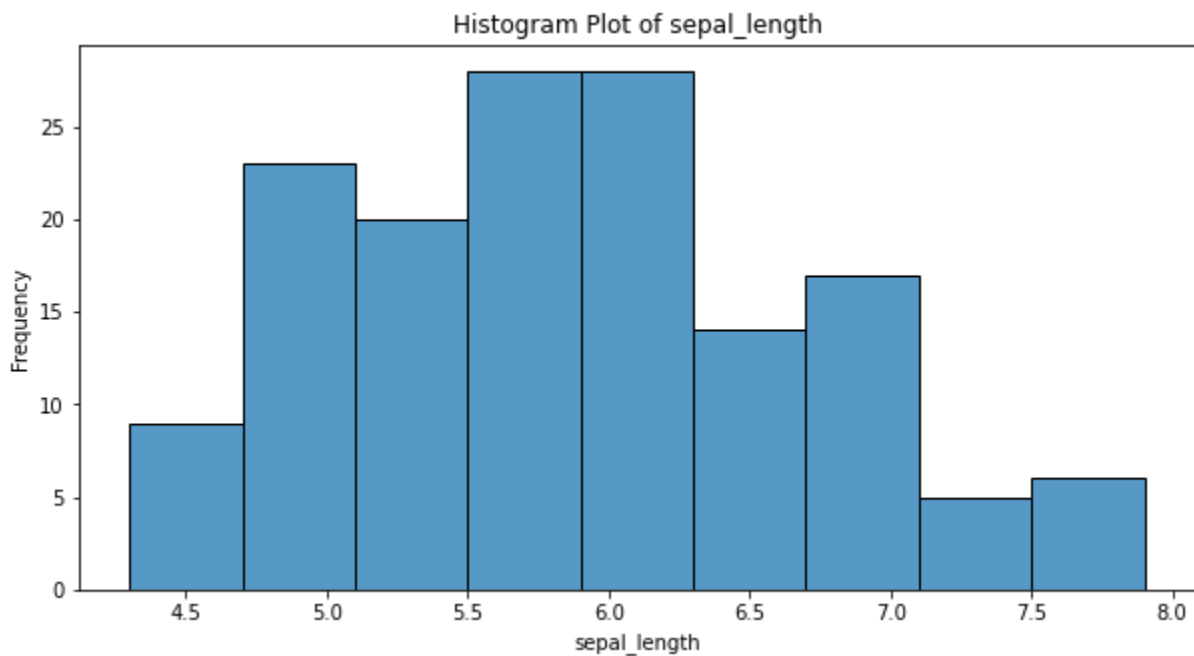
```


# Plotting histograms for each column
for col in iris.columns:
    plt.figure(figsize=(10, 5))
    sns.histplot(iris, x=col, kde=False)

```



```
plt.xlabel(col)
plt.ylabel('Frequency')
plt.title(f'Histogram Plot of {col}')
```





```
(data1.columns)
```

```
Index(['sepal length (cm)', 'sepal width (cm)', 'petal length (cm)',
      'petal width (cm)', 'target'],
      dtype='object')
```

```
#Normalization
```

```
data1['sepal length (cm)'] = data1['sepal length (cm)'] / (data1['sepal length (cm)'].max() - data1['sepal length (cm)'].min())
data1['sepal width (cm)'] = data1['sepal width (cm)'] / (data1['sepal width (cm)'].max() - data1['sepal width (cm)'].min())
data1['petal length (cm)'] = data1['petal length (cm)'] / (data1['petal length (cm)'].max() - data1['petal length (cm)'].min())
data1['petal width (cm)'] = data1['petal width (cm)'] / (data1['petal width (cm)'].max() - data1['petal width (cm)'].min())
```

```
data1
```

|     | sepal length<br>(cm) | sepal width<br>(cm) | petal length<br>(cm) | petal width<br>(cm) | target |
|-----|----------------------|---------------------|----------------------|---------------------|--------|
| 0   | 0.645570             | 0.795455            | 0.202899             | 0.08                | 0.0    |
| 1   | 0.620253             | 0.681818            | 0.202899             | 0.08                | 0.0    |
| 2   | 0.594937             | 0.727273            | 0.188406             | 0.08                | 0.0    |
| 3   | 0.582278             | 0.704545            | 0.217391             | 0.08                | 0.0    |
| 4   | 0.632911             | 0.818182            | 0.202899             | 0.08                | 0.0    |
| ... | ...                  | ...                 | ...                  | ...                 | ...    |
| 145 | 0.848101             | 0.681818            | 0.753623             | 0.92                | 2.0    |
| 146 | 0.797468             | 0.568182            | 0.724638             | 0.76                | 2.0    |
| 147 | 0.822785             | 0.681818            | 0.753623             | 0.80                | 2.0    |
| 148 | 0.784810             | 0.772727            | 0.782609             | 0.92                | 2.0    |
| 149 | 0.746835             | 0.681818            | 0.739130             | 0.72                | 2.0    |

```
150 rows x 5 columns
```

```
iris = datasets.load_iris()
```

```
X = iris.data
y = iris.target
```

```
# Split data into training and testing sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)
```

```

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)

# Calculate class probabilities
class_probs = np.zeros((3, X.shape[1]))
for i in range(3):
    class_probs[i,:] = np.mean(X_train[y_train==i], axis=0)

# Calculate class priors
class_priors = np.zeros(3)
for i in range(3):
    class_priors[i] = np.mean(y_train==i)

# Define predict function
def predict(X, class_probs, class_priors):
    probs = np.zeros((X.shape[0], 3))
    for i in range(3):
        probs[:,i] = class_priors[i] * np.prod(np.power(class_probs[i,:], X), axis=
    return np.argmax(probs, axis=1)

# Predict class labels on test data
y_pred = predict(X_test, class_probs, class_priors)

# Calculate accuracy
accuracy = np.mean(y_pred==y_test)
print("Accuracy:", accuracy)

    Accuracy: 0.8777777777777778

```

## Q4: Naive Bayes using SK-Learn

```

import numpy as np
from sklearn import datasets
from sklearn.model_selection import train_test_split
from sklearn.naive_bayes import GaussianNB

# Load Iris dataset
iris = datasets.load_iris()
X = iris.data
y = iris.target

# Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)

# Train Naive Bayes classifier
clf = GaussianNB()
clf.fit(X_train, y_train)

# Predict class labels on test data

```

```
y_pred = clf.predict(X_test)

# Calculate accuracy
accuracy = np.mean(y_pred==y_test)
print("Accuracy:", accuracy)

    Accuracy: 1.0
```

## Q5: KNN

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt

df = pd.read_csv("/content/drive/MyDrive/Datasets/bmd.csv") # Load the data

# Split the data into training and testing sets
X = df["age"].values.reshape(-1, 1)
y = df["bmd"].values
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2)


# Calculating the Euclidean distance between two points
def euclidean_distance(x1, x2):
    return np.sqrt(np.sum((x1 - x2)**2))

# Define a function to fit KNN Regressor
def knn_regressor(X_train, y_train, X_test, k):
    y_pred = []
    for x_test in X_test:
        distances = []
        for i, j in zip(X_train, y_train):
            distance = euclidean_distance(x_test, i)
            distances.append((distance, j))
        distances = sorted(distances, key=lambda x: x[0])
        k_neighbors = distances[:k]
        k_neighbors_y = [neighbor[1] for neighbor in k_neighbors]
        y_pred.append(np.mean(k_neighbors_y))
    return np.array(y_pred)

# Define a function to calculate R2 score
def r2_score(y_true, y_pred):
    mean_y_true = np.mean(y_true)
    sse = np.sum((y_true - y_pred)**2)
    sst = np.sum((y_true - mean_y_true)**2)
    return 1 - sse/sst
```

```
# Plot the results for different values of k
r2_scores = []
k_values = [1,2,3,4,5,6,7,8,9]
for k in k_values:
    y_pred = knn_regressor(X_train, y_train, X_test, k)
    r2 = r2_score(y_test, y_pred)
    r2_scores.append(r2)

plt.plot(k_values, r2_scores)
plt.xlabel("k")
plt.ylabel("R2 score")
plt.show()
```



```
ind = np.argmax(r2_scores, axis=None, out=None)
print("Best value of k: ", k_values[ind])
print("Corresponding R2 score: ", r2_scores[ind])
```

```
Best value of k: 8
Corresponding R2 score: -0.05781999422048245
```

[Colab paid products](#) - [Cancel contracts here](#)