

Q2 [P || CO2 & CO3] Autoencoder (6 points)

Pick random 1000 images (Train- 900, Test-100) of 5 Superclasses of the CIFAR-100 dataset ([https:// www.cs.toronto.edu/~kriz/cifar.html](https://www.cs.toronto.edu/~kriz/cifar.html), <https://www.cs.toronto.edu/~kriz/cifar-100-python.tar.gz>) and train an autoencoder to regenerate the images. Apply batch normalization and plot the loss vs. epoch training curve. Print a 5x2 grid containing 1 test image of each class, in which the first column contains the original image and the second column contains the autoencoder output of the same image. Now, use the latent embeddings learned in the above autoencoder to build a five-class classifier. Show performance on train and test sets using accuracy as an evaluation metric.

```
!pip install tensorflow
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Requirement already satisfied: tensorflow in /usr/local/lib/python3.9/dist-packages (2.11.0)
Requirement already satisfied: protobuf<3.20,>=3.9.2 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (3.19.6)
Requirement already satisfied: six>=1.12.0 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (1.16.0)
Requirement already satisfied: setuptools in /usr/local/lib/python3.9/dist-packages (from tensorflow) (63.4.3)
Requirement already satisfied: gast<0.4.0,>=0.2.1 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (0.4.0)
Requirement already satisfied: astunparse>=1.6.0 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (1.6.3)
Requirement already satisfied: grpcio<2.0,>=1.24.3 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (1.51.3)
Requirement already satisfied: h5py>=2.9.0 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (3.8.0)
Requirement already satisfied: numpy>=1.20 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (1.22.4)
Requirement already satisfied: tensorflow-estimator<2.12,>=2.11.0 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (2.11.0)
Requirement already satisfied: libclang>=13.0.0 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (15.0.6.1)
Requirement already satisfied: termcolor>=1.1.0 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (2.2.0)
Requirement already satisfied: tensorboard<2.12,>=2.11 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (2.11.2)
Requirement already satisfied: tensorflow-io-gcs-filesystem>=0.23.1 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (0.31.0)
Requirement already satisfied: typing-extensions>=3.6.6 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (4.5.0)
Requirement already satisfied: opt-einsum>=2.3.2 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (3.3.0)
Requirement already satisfied: absl-py>=1.0.0 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (1.4.0)
Requirement already satisfied: keras<2.12,>=2.11.0 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (2.11.0)
Requirement already satisfied: google-pasta>=0.1.1 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (0.2.0)
Requirement already satisfied: flatbuffers>=2.0 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (23.3.3)
Requirement already satisfied: packaging in /usr/local/lib/python3.9/dist-packages (from tensorflow) (23.0)
Requirement already satisfied: wrapt>=1.11.0 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (1.15.0)
Requirement already satisfied: wheel<1.0,>=0.23.0 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (0.40.0)
Requirement already satisfied: markdown>=2.6.8 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (3.4.3)
Requirement already satisfied: tensorboard-data-server<0.7.0,>=0.6.0 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (0.6.0)
Requirement already satisfied: werkzeug>=1.0.1 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (2.2.3)
Requirement already satisfied: google-auth<3,>=1.6.3 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (2.22.0)
Requirement already satisfied: requests<3,>=2.21.0 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (2.28.1)
Requirement already satisfied: tensorboard-plugin-wit>=1.6.0 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (1.8.0)
Requirement already satisfied: google-auth-oauthlib<0.5,>=0.4.1 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (0.4.6)
Requirement already satisfied: rsa<5,>=3.1.4 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (4.9)
Requirement already satisfied: cachetools<6.0,>=2.0.0 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (5.2.1)
Requirement already satisfied: pyasn1-modules>=0.2.1 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (0.3.0)
Requirement already satisfied: requests-oauthlib>=0.7.0 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (1.3.1)
Requirement already satisfied: importlib-metadata>=4.4 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (6.7.0)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (3.6)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (2024.7.4)
Requirement already satisfied: charset-normalizer<=2.0.0 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (3.3.0)
Requirement already satisfied: urllib3<1.27,>=1.21.1 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (2.2.2)
Requirement already satisfied: MarkupSafe>=2.1.1 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (2.1.5)
Requirement already satisfied: zipp>=0.5 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (3.17.0)
Requirement already satisfied: pyasn1<0.5.0,>=0.4.6 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (0.6.0)
Requirement already satisfied: oauthlib>=3.0.0 in /usr/local/lib/python3.9/dist-packages (from tensorflow) (3.2.2)
```

Double-click (or enter) to edit

```
import seaborn as sns
import numpy as np
import pandas as pd
import matplotlib
from matplotlib import pyplot as plt
from sklearn import svm, datasets
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score
from sklearn.model_selection import KFold
import pickle
import gzip
from google.colab import drive
import random
import keras
from keras.layers import Flatten,Dense,Conv2D,MaxPooling2D,UpSampling2D,Dropout
from keras.models import Model,load_model
from keras import Input,Sequential
from keras.layers import BatchNormalization
from keras.optimizers import RMSprop
from keras.utils import to_categorical
import tensorflow as tf
from keras.datasets import fashion_mnist
from keras.applications import VGG16
from keras.applications import VGG19
from keras.applications import ResNet50V2
from keras.applications import MobileNet
from keras.applications import EfficientNetB0
from keras.applications.mobilenet import preprocess_input
```

```

import cv2
import gc
from sklearn import preprocessing

tf.config.run_functions_eagerly(True)

drive.mount('/content/gdrive')

Mounted at /content/gdrive

!tar -xvf "/content/gdrive/MyDrive/ML_A3/cifar-100-python.tar.gz"

cifar-100-python/
cifar-100-python/file.txt~
cifar-100-python/train
cifar-100-python/test
cifar-100-python/meta

with open("/content/cifar-100-python/train", 'rb') as data:
    train = pickle.load(data, encoding='bytes')

with open("/content/cifar-100-python/test", 'rb') as data:
    test = pickle.load(data, encoding='bytes')

with open("/content/cifar-100-python/meta", 'rb') as data:
    meta = pickle.load(data, encoding='bytes')

print(train.keys())
print(len(train[b'data']))
print(train[b'data'][0])
print(len(train[b'data'][0]))

dict_keys([b'filenames', b'batch_label', b'fine_labels', b'coarse_labels', b'data'])
50000
[255 255 255 ... 10 59 79]
3072

superclasses = np.unique(np.concatenate((np.array(train[b'coarse_labels']), np.array(test[b'coarse_labels'])), axis=0))
print(superclasses)

[ 0  1  2  3  4  5  6  7  8  9 10 11 12 13 14 15 16 17 18 19]

random.seed(12223)

selected_classes = random.sample(list(superclasses), 5)
print(selected_classes)

[2, 8, 10, 18, 0]

dict_of_indices = {}
for i in selected_classes:
    dict_of_indices[i]=[]
    for j in range(len(train[b'coarse_labels'])):
        if train[b'coarse_labels'][j]==i:
            dict_of_indices[i].append(j)
print(dict_of_indices)

{2: [12, 22, 70, 126, 141, 164, 167, 174, 221, 225, 246, 274, 314, 318, 324, 330, 336, 352, 362, 419, 445, 456, 460, 466, 472, 498, 51

x_train = []
y_train = np.array([])
x_indices_present = []
while len(x_train)!=900:
    classChosen=selected_classes[random.randrange(len(selected_classes))]
    elementChosen = dict_of_indices[classChosen][random.randrange(len(dict_of_indices[classChosen]))]
    if elementChosen not in x_indices_present:
        y_train=np.concatenate((y_train,np.array([classChosen])),axis=0)
        x_train.append(np.array(train[b'data'][elementChosen]))
        x_indices_present.append(elementChosen)
x_train = np.array(x_train)
dict_of_indices = {}
for i in selected_classes:
    dict_of_indices[i]=[]
    for j in range(len(test[b'coarse_labels'])):

```

```

        if test[b'coarse_labels'][j]==i:
            dict_of_indices[i].append(j)
x_test = []
y_test = np.array([])
x_indices_present = []
while len(x_test)!=100:
    classChosen=selected_classes[random.randrange(len(selected_classes))]
    elementChosen = dict_of_indices[classChosen][random.randrange(len(dict_of_indices[classChosen]))]
    if elementChosen not in x_indices_present:
        y_test=np.concatenate((y_test,np.array([classChosen])),axis=0)
        x_test.append(np.array(test[b'data'][elementChosen]))
        x_indices_present.append(elementChosen)
x_test = np.array(x_test)

print(type(x_train[0]))
print(len(x_train))
print(len(y_train))
print(len(x_test))
print(len(y_test))

<class 'numpy.ndarray'>
900
900
100
100

x_train = x_train / np.max(x_train)
x_test = x_test / np.max(x_test)

x_train = x_train.reshape(-1,3,32,32).transpose(0, 2, 3, 1)
x_test = x_test.reshape(-1,3,32,32).transpose(0, 2, 3, 1)

print(x_train.shape)
print(x_test.shape)

(900, 32, 32, 3)
(100, 32, 32, 3)

from tensorflow.keras.layers import Conv2D, BatchNormalization, MaxPooling2D

def encoder(input_tensor):
    # First convolutional layer
    conv1 = Conv2D(filters=32, # 32 output channels
                   kernel_size=(3,3), # 3x3 kernel size
                   activation='relu', # ReLU activation function
                   padding='same')(input_tensor) # Same padding
    # Batch normalization layer after conv1
    bn1 = BatchNormalization()(conv1)
    # Max pooling layer after bn1
    maxpool1 = MaxPooling2D(pool_size=(2,2), padding='same')(bn1)

    # Second convolutional layer
    conv2 = Conv2D(filters=64, kernel_size=(3,3), activation='relu', padding='same')(maxpool1)
    # Batch normalization layer after conv2
    bn2 = BatchNormalization()(conv2)

    # Third convolutional layer
    conv3 = Conv2D(filters=128, kernel_size=(3,3), activation='relu', padding='same')(bn2)
    # Batch normalization layer after conv3
    bn3 = BatchNormalization()(conv3)
    # Max pooling layer after bn3
    maxpool2 = MaxPooling2D(pool_size=(2,2), padding='same')(bn3)

    # Fourth convolutional layer
    conv4 = Conv2D(filters=256, kernel_size=(3,3), activation='relu', padding='same')(maxpool2)
    # Batch normalization layer after conv4
    bn4 = BatchNormalization()(conv4)

    # Final convolutional layer to encode the input image
    encoded_img = Conv2D(filters=16, kernel_size=(3,3), activation='sigmoid', padding='same')(conv4_layer)
    # Use sigmoid activation for the last layer to squash the pixel values between 0 and 1, making the output a binary image

    # Return the encoded image
    return encoded_img

def my_decoder(encoded_img):
    # Decoder network
    conv1_layer = Conv2D(filters=16, kernel_size=(3,3), activation='relu', padding='same')(encoded_img)

```

```

conv1_layer = BatchNormalization()(conv1_layer)

conv2_layer = Conv2D(filters=32, kernel_size=(3,3), activation='relu', padding='same')(conv1_layer)
conv2_layer = BatchNormalization()(conv2_layer)

# Upsampling layer 1
upsampled_layer1 = UpSampling2D(size=(2,2))(conv2_layer)

conv3_layer = Conv2D(filters=64, kernel_size=(3,3), activation='relu', padding='same')(upsampled_layer1)
conv3_layer = BatchNormalization()(conv3_layer)

conv4_layer = Conv2D(filters=128, kernel_size=(3,3), activation='relu', padding='same')(conv3_layer)
conv4_layer = BatchNormalization()(conv4_layer)

# Upsampling layer 2
upsampled_layer2 = UpSampling2D(size=(2,2))(conv4_layer)

decoded_img = Conv2D(filters=3, kernel_size=(3,3), activation='sigmoid', padding='same')(upsampled_layer2)
# Use sigmoid activation for the last layer to squash the pixel values between 0 and 1, making the output a binary image

# Return the decoded image
return decoded_img

```

```

autoEnc = my_decoder(encoder())
print(autoEnc.summary())
autoEnc.compile(optimizer = RMSprop(),loss="mean_squared_error",metrics=['mean_squared_error'])

```

max_pooling2d (MaxPooling2D)	(None, 16, 16, 64)	0
conv2d_3 (Conv2D)	(None, 16, 16, 64)	36928
batch_normalization_3 (Batch Normalization)	(None, 16, 16, 64)	256
conv2d_4 (Conv2D)	(None, 16, 16, 32)	18464
batch_normalization_4 (Batch Normalization)	(None, 16, 16, 32)	128
conv2d_5 (Conv2D)	(None, 16, 16, 16)	4624
batch_normalization_5 (Batch Normalization)	(None, 16, 16, 16)	64
conv2d_6 (Conv2D)	(None, 16, 16, 16)	2320
batch_normalization_6 (Batch Normalization)	(None, 16, 16, 16)	64
conv2d_7 (Conv2D)	(None, 16, 16, 32)	4640
batch_normalization_7 (Batch Normalization)	(None, 16, 16, 32)	128
conv2d_8 (Conv2D)	(None, 16, 16, 64)	18496
up_sampling2d (UpSampling2D)	(None, 32, 32, 64)	0
batch_normalization_8 (Batch Normalization)	(None, 32, 32, 64)	256
conv2d_9 (Conv2D)	(None, 32, 32, 32)	18464
batch_normalization_9 (Batch Normalization)	(None, 32, 32, 32)	128
conv2d_10 (Conv2D)	(None, 32, 32, 16)	4624
batch_normalization_10 (Batch Normalization)	(None, 32, 32, 16)	64
conv2d_11 (Conv2D)	(None, 32, 32, 3)	435
batch_normalization_11 (Batch Normalization)	(None, 32, 32, 3)	12

```

=====
Total params: 134,127
Trainable params: 133,353
Non-trainable params: 774

```

None

```

trackLoss = autoEnc.fit(x_train,x_train,batch_size=100,epochs=50,validation_split=0.3,callbacks = [
    keras.callbacks.ModelCheckpoint("autoencoder.h5",save_best_only=True),
    keras.callbacks.EarlyStopping(monitor="val_loss", patience=25, verbose=1),
])

Epoch 4/50
7/7 [=====] - 1s 159ms/step - loss: 0.7525 - mean_squared_error: 0.7525 - val_loss: 0.0564 - val_mean_squared_error: 0.0564
Epoch 5/50
7/7 [=====] - 1s 151ms/step - loss: 0.7227 - mean_squared_error: 0.7227 - val_loss: 0.0560 - val_mean_squared_error: 0.0560
Epoch 6/50
7/7 [=====] - 1s 131ms/step - loss: 0.7004 - mean_squared_error: 0.7004 - val_loss: 0.0574 - val_mean_squared_error: 0.0574
Epoch 7/50
7/7 [=====] - 1s 156ms/step - loss: 0.6799 - mean_squared_error: 0.6799 - val_loss: 0.0530 - val_mean_squared_error: 0.0530
Epoch 8/50
7/7 [=====] - 1s 133ms/step - loss: 0.6615 - mean_squared_error: 0.6615 - val_loss: 0.0605 - val_mean_squared_error: 0.0605
Epoch 9/50
7/7 [=====] - 1s 131ms/step - loss: 0.6358 - mean_squared_error: 0.6358 - val_loss: 0.0573 - val_mean_squared_error: 0.0573
Epoch 10/50
7/7 [=====] - 1s 143ms/step - loss: 0.6024 - mean_squared_error: 0.6024 - val_loss: 0.0594 - val_mean_squared_error: 0.0594
Epoch 11/50
7/7 [=====] - 1s 167ms/step - loss: 0.5389 - mean_squared_error: 0.5389 - val_loss: 0.0605 - val_mean_squared_error: 0.0605
Epoch 12/50
7/7 [=====] - 1s 184ms/step - loss: 0.4597 - mean_squared_error: 0.4597 - val_loss: 0.0606 - val_mean_squared_error: 0.0606
Epoch 13/50
7/7 [=====] - 1s 183ms/step - loss: 0.4218 - mean_squared_error: 0.4218 - val_loss: 0.0618 - val_mean_squared_error: 0.0618
Epoch 14/50
7/7 [=====] - 1s 130ms/step - loss: 0.3926 - mean_squared_error: 0.3926 - val_loss: 0.0605 - val_mean_squared_error: 0.0605
Epoch 15/50
7/7 [=====] - 1s 130ms/step - loss: 0.3692 - mean_squared_error: 0.3692 - val_loss: 0.0675 - val_mean_squared_error: 0.0675
Epoch 16/50
7/7 [=====] - 1s 132ms/step - loss: 0.3520 - mean_squared_error: 0.3520 - val_loss: 0.0664 - val_mean_squared_error: 0.0664
Epoch 17/50
7/7 [=====] - 1s 147ms/step - loss: 0.3563 - mean_squared_error: 0.3563 - val_loss: 0.0603 - val_mean_squared_error: 0.0603
Epoch 18/50
7/7 [=====] - 1s 142ms/step - loss: 0.3265 - mean_squared_error: 0.3265 - val_loss: 0.0637 - val_mean_squared_error: 0.0637
Epoch 19/50
7/7 [=====] - 1s 142ms/step - loss: 0.3159 - mean_squared_error: 0.3159 - val_loss: 0.0612 - val_mean_squared_error: 0.0612
Epoch 20/50
7/7 [=====] - 1s 131ms/step - loss: 0.3038 - mean_squared_error: 0.3038 - val_loss: 0.0658 - val_mean_squared_error: 0.0658
Epoch 21/50
7/7 [=====] - 1s 130ms/step - loss: 0.2926 - mean_squared_error: 0.2926 - val_loss: 0.0658 - val_mean_squared_error: 0.0658
Epoch 22/50
7/7 [=====] - 1s 132ms/step - loss: 0.2818 - mean_squared_error: 0.2818 - val_loss: 0.0621 - val_mean_squared_error: 0.0621
Epoch 23/50
7/7 [=====] - 1s 131ms/step - loss: 0.2722 - mean_squared_error: 0.2722 - val_loss: 0.0785 - val_mean_squared_error: 0.0785
Epoch 24/50
7/7 [=====] - 1s 155ms/step - loss: 0.2618 - mean_squared_error: 0.2618 - val_loss: 0.0687 - val_mean_squared_error: 0.0687
Epoch 25/50
7/7 [=====] - 1s 174ms/step - loss: 0.2547 - mean_squared_error: 0.2547 - val_loss: 0.0691 - val_mean_squared_error: 0.0691
Epoch 26/50
7/7 [=====] - 1s 190ms/step - loss: 0.2470 - mean_squared_error: 0.2470 - val_loss: 0.0634 - val_mean_squared_error: 0.0634
Epoch 27/50
7/7 [=====] - 1s 136ms/step - loss: 0.2481 - mean_squared_error: 0.2481 - val_loss: 0.0625 - val_mean_squared_error: 0.0625
Epoch 28/50
7/7 [=====] - 1s 143ms/step - loss: 0.2382 - mean_squared_error: 0.2382 - val_loss: 0.0624 - val_mean_squared_error: 0.0624
Epoch 29/50
7/7 [=====] - 1s 142ms/step - loss: 0.2202 - mean_squared_error: 0.2202 - val_loss: 0.0641 - val_mean_squared_error: 0.0641
Epoch 30/50
7/7 [=====] - 1s 142ms/step - loss: 0.2204 - mean_squared_error: 0.2204 - val_loss: 0.0709 - val_mean_squared_error: 0.0709
Epoch 31/50
7/7 [=====] - 1s 130ms/step - loss: 0.2083 - mean_squared_error: 0.2083 - val_loss: 0.0709 - val_mean_squared_error: 0.0709
Epoch 32/50
7/7 [=====] - 1s 143ms/step - loss: 0.2045 - mean_squared_error: 0.2045 - val_loss: 0.0666 - val_mean_squared_error: 0.0666
Epoch 32: early stopping

```

```

with open('autoencoder.pkl','wb') as data:
    pickle.dump(trackLoss.history,data)

```

```

with open('autoencoder.pkl','rb') as data:
    trackLoss = pickle.load(data)

```

```

plt.plot(trackLoss['val_loss'])
plt.plot(trackLoss['loss'])
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend(['Validation Loss', 'Train Loss'])
plt.show()

```



```
autoEnc = load_model("autoencoder.h5")
```

```
len(np.unique(y_train))
```

```
5
```


```
indexes_taken = {}
for i in selected_classes:
    indexes_taken[i]=[]
    for j in range(len(y_test)):
        if y_test[j]==i:
            indexes_taken[i].append(j)
print(indexes_taken)
```

```
{2: [0, 13, 19, 24, 26, 28, 32, 34, 43, 46, 48, 58, 60, 67, 69, 70, 71, 73, 82, 92, 97], 8: [10, 14, 15, 17, 18, 22, 31, 37, 44, 57, 58, 61, 62, 63, 64, 65, 66, 68, 72, 74, 75, 76, 77, 78, 79, 80, 81, 83, 84, 85, 86, 87, 88, 89, 90, 91, 93, 94, 95, 96, 98, 99]}
```

```
y_pred = autoEnc.predict(x_test)
```

```
4/4 [=====] - 0s 64ms/step
```

```
y_pred = y_pred/np.max(y_pred)
for j in selected_classes:
    fig, ax = plt.subplots(1,2)
    indexChosen = indexes_taken[j][random.randrange(len(indexes_taken[j]))]
    ax[0].imshow(x_test[indexChosen])
    ax[1].imshow(y_pred[indexChosen])
    plt.show()
```




```
def classifier(cl):
    cl.add(Flatten())
    cl.add(Dense(128, activation='relu'))
    cl.add(Dense(64, activation='relu'))
    cl.add(Dense(5, activation='softmax'))
    return cl
```

```
classifier = classifier(encoding())
```

```
print(classifier.summary())
```

```
Model: "sequential_17"
```


Layer (type)	Output Shape	Param #
=====		
conv2d_12 (Conv2D)	(None, 32, 32, 16)	448
batch_normalization_12 (Batch Normalization)	(None, 32, 32, 16)	64
conv2d_13 (Conv2D)	(None, 32, 32, 32)	4640
batch_normalization_13 (Batch Normalization)	(None, 32, 32, 32)	128
conv2d_14 (Conv2D)	(None, 32, 32, 64)	18496
batch_normalization_14 (Batch Normalization)	(None, 32, 32, 64)	256
max_pooling2d_4 (MaxPooling2D)	(None, 16, 16, 64)	0
conv2d_15 (Conv2D)	(None, 16, 16, 64)	36928
batch_normalization_15 (Batch Normalization)	(None, 16, 16, 64)	256
conv2d_16 (Conv2D)	(None, 16, 16, 32)	18464
batch_normalization_16 (Batch Normalization)	(None, 16, 16, 32)	128
conv2d_17 (Conv2D)	(None, 16, 16, 16)	4624
batch_normalization_17 (Batch Normalization)	(None, 16, 16, 16)	64
flatten_5 (Flatten)	(None, 4096)	0
dense_15 (Dense)	(None, 128)	524416
dense_16 (Dense)	(None, 64)	8256
dense_17 (Dense)	(None, 5)	325
=====		
Total params: 617,493		
Trainable params: 617,045		
Non-trainable params: 448		
None		

```
for i in range(13):
    classifier.layers[i].set_weights(autoEnc.layers[i].get_weights())
    classifier.layers[i].trainable = False
```

```
classifier.compile(optimizer = 'adam', loss="categorical_crossentropy", metrics=['accuracy'])
```

```
with open('classifier.pkl', 'rb') as data:
    trackLossClassifier = pickle.load(data)
```

```
plt.plot(trackLossClassifier['val_loss'])
plt.plot(trackLossClassifier['loss'])
plt.legend(['Validation Loss', 'Train Loss'])
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.show()
```

```

classifier = load_model("Classifier.h5")

print("Training accuracy is "+str((classifier.evaluate(x_train,y_train_labelled))[1]*100))
print("Test accuracy is "+str((classifier.evaluate(x_test,y_test_labelled))[1]*100))

29/29 [=====] - 2s 51ms/step - loss: 1.0244 - accuracy: 0.6056
Training accuracy is 60.5555534362793
4/4 [=====] - 0s 62ms/step - loss: 1.1475 - accuracy: 0.6000
Test accuracy is 60.00000238418579

```

Q3 [P || CO3 & C4] SVM (10 points)

Use the following SVM strategies to classify sapodillas from kiwi for the dataset available here¹ . For each strategy, plot the learned decision boundary. Use an appropriate evaluation metric to present the test scores. (a) Linear SVM (b) Polynomial SVM (c) Kernel SVM Finally, write the individual and combined inferences of the obtained results.

```

df = pd.read_csv('/content/gdrive/MyDrive/ML_A3/Sapodillas_and_Kiwis_MLA3.csv')

x = df[['Weight', 'Size']]

y = df['Class']

x_train, x_test, y_train, y_test = train_test_split(x,y,test_size=0.3,random_state=15)
x_of_df = x.copy()
y_of_df = y.copy()

# normalising the data
for col in x_train.columns:
    if col=='Class':
        continue
    # print(col)
    x_test[col]=(x_test[col] - x_train[col].min()) / (x_train[col].max()- x_train[col].min())
    x_of_df[col]=(x_of_df[col] - x_train[col].min()) / (x_of_df[col].max()- x_train[col].min())
    x_train[col]=(x_train[col] - x_train[col].min()) / (x_train[col].max()- x_train[col].min())

# normalising the data based on the training data

print(x_train)
# printing the training data
print(y_train)

```

	Weight	Size
1	0.4	0.112360
2	0.0	0.044944
39	0.8	0.573034
6	0.5	0.870787
26	0.2	0.000000
34	0.3	0.460674
38	0.5	0.887640
37	0.9	0.825843
33	0.8	0.651685
9	0.0	0.146067
4	0.2	0.387640
29	0.6	0.702247
10	0.8	1.000000
36	0.4	0.421348
23	0.3	0.039326
13	0.3	0.258427
17	1.0	0.696629
15	0.0	0.264045
22	0.4	0.337079
21	0.5	0.460674
11	0.5	0.820225
7	1.0	0.617978
27	0.5	0.117978
28	0.9	0.696629
0	0.4	0.213483
5	0.8	0.938202

```

12 0.9 0.853933
8 0.9 0.758427
1 Sapodilla
2 Sapodilla
39 Kiwi
6 Kiwi
26 Sapodilla
34 Sapodilla
38 Kiwi
37 Kiwi
33 Kiwi
9 Sapodilla
4 Sapodilla
29 Kiwi
10 Kiwi
36 Sapodilla
23 Sapodilla
13 Sapodilla
17 Kiwi
15 Sapodilla
22 Sapodilla
21 Sapodilla
11 Kiwi
7 Kiwi
27 Sapodilla
28 Kiwi
0 Sapodilla
5 Kiwi
12 Kiwi
8 Kiwi
Name: Class, dtype: object

```

```

classifier = svm.SVC(kernel='linear', C = 1).fit(x_train, y_train)

```

```

y_pred = classifier.predict(x_test)
linear_acc_score = accuracy_score(y_test, y_pred)
print("Accuracy score for linear SVM "+str(linear_acc_score))

```

```

Accuracy score for linear SVM 1.0

```

```

min_wt, max_wt = x_of_df['Weight'].min() - 1, x_of_df['Weight'].max() + 1
min_size, max_size = x_of_df['Size'].min() - 1, x_of_df['Size'].max() + 1
x_axis_val, y_axis_val = np.meshgrid(np.arange(min_wt, max_wt, 0.005), np.arange(min_size, max_size, 0.005))
boundary = classifier.predict(np.c_[x_axis_val.ravel(), y_axis_val.ravel()])

```

```

boundary = preprocessing.LabelEncoder().fit_transform(boundary)
boundary = boundary.reshape(x_axis_val.shape)
plt.contour(x_axis_val, y_axis_val, boundary)
plt.xlabel('Weight')
plt.ylabel('Size')
plt.title('Linear SVM ')
sns.scatterplot(x=x_of_df['Weight'], y=x_of_df['Size'], hue=y)

```

```


plt.show()

```

```

/usr/local/lib/python3.9/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but SVC was fitted with
warnings.warn(

```



```

print(np.array(y_pred))
print(np.array(y_test))

```

```

['Kiwi' 'Sapodilla' 'Sapodilla' 'Sapodilla' 'Sapodilla' 'Kiwi' 'Kiwi'
 'Kiwi' 'Sapodilla' 'Kiwi' 'Kiwi' 'Sapodilla']
['Kiwi' 'Sapodilla' 'Sapodilla' 'Sapodilla' 'Sapodilla' 'Kiwi' 'Kiwi'
 'Kiwi' 'Sapodilla' 'Kiwi' 'Kiwi' 'Sapodilla']

```

```

from sklearn.model_selection import KFold
from sklearn import svm
from sklearn.metrics import accuracy_score
import numpy as np

```

```

# Set up a KFold cross-validation object with 5 splits
kFold = KFold(n_splits=5)

```

```

kFold = KFold(n_splits=5)

# Initialize an empty list to store the accuracy scores for each fold
accuracy_scores = []

# Iterate through each fold
for i, (train_index, test_index) in enumerate(kFold.split(x_of_df)):

    # Initialize empty lists to store the training and testing data for this fold
    x_train_k = []
    y_train_k = []
    x_test_k = []
    y_test_k = []

    # Retrieve the training and testing data for this fold from the input data
    for j in train_index:
        x_train_k.append(np.array(x_of_df.loc[j]))
        y_train_k.append(np.array(y_of_df.loc[j]))
    for j in test_index:
        x_test_k.append(np.array(x_of_df.loc[j]))
        y_test_k.append(np.array(y_of_df.loc[j]))

    # Convert the training and testing data into numpy arrays
    x_train_k = np.array(x_train_k)
    y_train_k = np.array(y_train_k)

    # Train a linear SVM model on the training data for this fold
    classifier = svm.SVC(kernel='linear', C=1).fit(x_train_k, y_train_k)

    # Use the trained model to predict the labels for the testing data for

    y_pred_k = classifier.predict(x_test_k)
    accuracy_scores.append(accuracy_score(y_test_k, y_pred_k))

print("5- fold cross validation accuracy score of linear SVM model is "+str(sum(accuracy_scores)/len(accuracy_scores)))

5- fold cross validation accuracy score of linear SVM model is 1.0

classifier = svm.SVC(kernel='poly', C = 1).fit(x_train, y_train)

y_pred = classifier.predict(x_test)
polynomial_acc_score = accuracy_score(y_test, y_pred)
print("Accuracy score for polynomial SVM "+str(polynomial_acc_score))

Accuracy score for polynomial SVM 1.0

kFold = KFold(n_splits=5)

accuracy_scores=[]
for i, (train_index, test_index) in enumerate(kFold.split(x_of_df)):
    x_train_k= []
    y_train_k = []
    x_test_k= []
    y_test_k = []
    for j in train_index:

        x_train_k.append(np.array(x_of_df.loc[j]))

        y_train_k.append(np.array(y_of_df.loc[j]))
    for j in test_index:
        x_test_k.append(np.array(x_of_df.loc[j]))

        y_test_k.append(np.array(y_of_df.loc[j]))
    x_train_k = np.array(x_train_k)
    y_train_k = np.array(y_train_k)
    classifier = svm.SVC(kernel='poly', C = 1).fit(x_train_k, y_train_k)

    y_pred_k = classifier.predict(x_test_k)
    accuracy_scores.append(accuracy_score(y_test_k, y_pred_k))

print("5- fold cross validation accuracy score of polynomial SVM model is "+str(sum(accuracy_scores)/len(accuracy_scores)))

5- fold cross validation accuracy score of polynomial SVM model is 1.0

min_wt, max_wt = x_of_df['Weight'].min() - 1, x_of_df['Weight'].max() + 1
min_size, max_size = x_of_df['Size'].min() - 1, x_of_df['Size'].max() + 1
x_axis_val, y_axis_val = np.meshgrid(np.arange(min_wt, max_wt, 0.005), np.arange(min_size, max_size, 0.005))
boundary = classifier.predict(np.c_[x_axis_val.ravel(), y_axis_val.ravel()])

boundary = preprocessing.LabelEncoder().fit_transform(boundary)

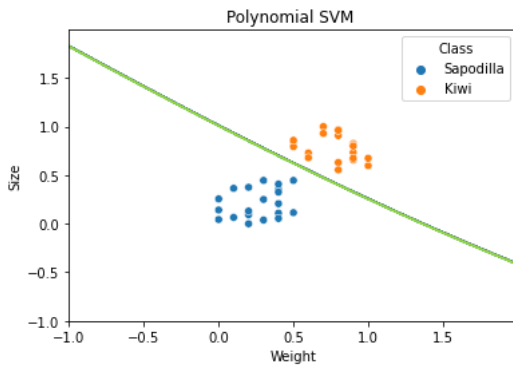
```

```

boundary = boundary.reshape(x_axis_val.shape)
plt.contour(x_axis_val, y_axis_val, boundary)
plt.xlabel('Weight')
plt.ylabel('Size')
plt.title('Polynomial SVM ')
sns.scatterplot(x=x_of_df['Weight'], y=x_of_df['Size'], hue=y)

```

```
plt.show()
```



```
classifier = svm.SVC(kernel='rbf', C = 1).fit(x_train, y_train)
```

```

y_pred = classifier.predict(x_test)
kernel_acc_score = accuracy_score(y_test, y_pred)
print("Accuracy score for ref kernel SVM "+str(kernel_acc_score))

```

```
Accuracy score for ref kernel SVM 1.0
```

```
kFold = KFold(n_splits=5)
```

```

accuracy_scores=[]
for i, (train_index, test_index) in enumerate(kFold.split(x_of_df)):
    x_train_k= []
    y_train_k = []
    x_test_k= []
    y_test_k = []
    for j in train_index:
        x_train_k.append(np.array(x_of_df.loc[j]))
        y_train_k.append(np.array(y_of_df.loc[j]))
    for j in test_index:
        x_test_k.append(np.array(x_of_df.loc[j]))

        y_test_k.append(np.array(y_of_df.loc[j]))
    x_train_k = np.array(x_train_k)
    y_train_k = np.array(y_train_k)
    classifier = svm.SVC(kernel='rbf', C = 1).fit(x_train_k, y_train_k)

    y_pred_k = classifier.predict(x_test_k)
    accuracy_scores.append(accuracy_score(y_test_k, y_pred_k))

```

```
print("5- fold cross validation accuracy score of polynomial SVM model is "+str(sum(accuracy_scores)/len(accuracy_scores)))
```

```
5- fold cross validation accuracy score of polynomial SVM model is 1.0
```

```

min_wt, max_wt = x_of_df['Weight'].min() - 1, x_of_df['Weight'].max() + 1
min_size, max_size = x_of_df['Size'].min() - 1, x_of_df['Size'].max() + 1
x_axis_val, y_axis_val = np.meshgrid(np.arange(min_wt, max_wt, 0.005), np.arange(min_size, max_size, 0.005))
boundary = classifier.predict(np.c_[x_axis_val.ravel(), y_axis_val.ravel()])


```

```

boundary = preprocessing.LabelEncoder().fit_transform(boundary)
boundary = boundary.reshape(x_axis_val.shape)
plt.contour(x_axis_val, y_axis_val, boundary)
plt.xlabel('Weight')
plt.ylabel('Size')
plt.title('Kernel SVM ')
sns.scatterplot(x=x_of_df['Weight'], y=x_of_df['Size'], hue=y)

```

```
plt.show()
```



```
print("The accuracy for kernel SVM "+str(kernel_acc_score))
print("The accuracy for polynomial SVM "+str(polynomial_acc_score))
print("The accuracy for linear SVM is "+str(linear_acc_score))
```

```
The accuracy for kernel SVM 1.0
The accuracy for polynomial SVM 1.0
The accuracy for linear SVM is 1.0
```

Q4 [P || CO3 & C4] Transfer Learning (20 points)

Use the FashionMNIST dataset (https://keras.io/api/datasets/fashion_mnist/) from Keras to build one classifier each based on the following pre-trained architectures trained on IMAGENET dataset (reference: <https://keras.io/api/applications/>). Remove the last layer, add your own few dense layers and the output layer. Train only these added layers from scratch on the data. Use an appropriate evaluation metric.

```
(x_train,y_train),(x_test,y_test)= fashion_mnist.load_data()
```

```
print(x_train.shape)
```

```
(60000, 28, 28)
```

```
print(x_train.shape)
```

```
(60000, 28, 28)
```

```
x_train_final=[]
x_test_final=[]
for i in x_train:
    x_train_final.append(cv2.resize(i, (56, 56)))
for i in x_test:
    x_test_final.append(cv2.resize(i, (56, 56)))
x_train_final = np.array(x_train_final)
x_test_final = np.array(x_test_final)
```

```
print(x_train_final.shape)
```

```
(60000, 56, 56)
```

```
x_train_final = np.dstack([x_train_final] * 3)
x_test_final = np.dstack([x_test_final] * 3)
```

```
x_train_final = x_train_final.reshape(-1,56,56,3)
x_test_final = x_test_final.reshape(-1,56,56,3)
```

```
print(x_train_final.shape)
```

```
(60000, 56, 56, 3)
```

```
print(np.unique(y_train))
print(np.unique(y_test))
```

```
[0 1 2 3 4 5 6 7 8 9]
[0 1 2 3 4 5 6 7 8 9]
```

```
y_train_labelled = to_categorical(y_train)
y_test_labelled = to_categorical(y_test)
```

VGG16

```

modVgg16 = Sequential()
modVgg16.add(VGG16(include_top=False,weights='imagenet',input_shape=(56,56, 3)))

```

```

for i in range(len(modVgg16.layers)):
    modVgg16.layers[i].trainable = False
print(modVgg16.summary())

```

```

Model: "sequential_18"

```

Layer (type)	Output Shape	Param #
vgg16 (Functional)	(None, 1, 1, 512)	14714688

```

=====
Total params: 14,714,688
Trainable params: 0
Non-trainable params: 14,714,688
None

```

```

modVgg16.add(Flatten())
modVgg16.add(Dense(128, activation='relu'))
modVgg16.add(Dense(64, activation='relu'))
modVgg16.add(Dense(10, activation='softmax'))
print(modVgg16.summary())

```

```

Model: "sequential_18"

```

Layer (type)	Output Shape	Param #
vgg16 (Functional)	(None, 1, 1, 512)	14714688
flatten_6 (Flatten)	(None, 512)	0
dense_18 (Dense)	(None, 128)	65664
dense_19 (Dense)	(None, 64)	8256
dense_20 (Dense)	(None, 10)	650

```

=====
Total params: 14,789,258
Trainable params: 74,570
Non-trainable params: 14,714,688
None

```

```


modVgg16.compile(optimizer = 'adam',loss="categorical_crossentropy",metrics=['accuracy'])

```

```

with open('VGG16classifier.pkl','rb') as data:
    trackLossClassifierVgg16 = pickle.load(data)
plt.plot(trackLossClassifierVgg16['val_loss'])
plt.plot(trackLossClassifierVgg16['loss'])
plt.legend(['Validation Loss', 'Train Loss'])
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.show()

```



```

modVgg16 = load_model("VGG16Classifier.h5")
val_vgg16 = (modVgg16.evaluate(x_test_final,y_test_labelled))[1]*100
print("Training accuracy is "+str((modVgg16.evaluate(x_train_final,y_train_labelled))[1]*100))
print("Test accuracy is "+str((modVgg16.evaluate(x_test_final,y_test_labelled))[1]*100))

```

```

/usr/local/lib/python3.9/dist-packages/tensorflow/python/data/ops/structured_function.py:256: UserWarning: Even though the `tf.config.
warnings.warn(
313/313 [=====] - 15s 47ms/step - loss: 0.4988 - accuracy: 0.8281
1875/1875 [=====] - 55s 29ms/step - loss: 0.3874 - accuracy: 0.8611
Training accuracy is 86.1050009727478

```

```
313/313 [=====] - 9s 28ms/step - loss: 0.4988 - accuracy: 0.8281
Test accuracy is 82.81000256538391
```

```
gc.collect()

37843
```

▼ VGG19

```
# Initialize a sequential model in Keras
modVgg19 = Sequential()

# Add the VGG19 model with pretrained weights from ImageNet, and freeze all of its layers
modVgg19.add(VGG19(include_top=False,weights='imagenet',input_shape=(56,56, 3)))
for i in range(len(modVgg19.layers)):
    modVgg19.layers[i].trainable = False

# Print a summary of the model architecture
print(modVgg19.summary())

# Add a flattening layer to convert the output of the VGG19 model into a 1D array
modVgg19.add(Flatten())

# Add a fully connected layer with 128 neurons and ReLU activation
modVgg19.add(Dense(128, activation='relu'))

# Add a second fully connected layer with 64 neurons and ReLU activation
modVgg19.add(Dense(64, activation='relu'))

# Add a final fully connected layer with 10 neurons and softmax activation
modVgg19.add(Dense(10, activation='softmax'))

# Print a summary of the model architecture
print(modVgg19.summary())

# Compile the model using the Adam optimizer, categorical crossentropy loss, and accuracy as a metric
modVgg19.compile(optimizer = 'adam',loss="categorical_crossentropy",metrics=[ 'accuracy' ])
```

Model: "sequential_19"

Layer (type)	Output Shape	Param #
vgg19 (Functional)	(None, 1, 1, 512)	20024384

Total params: 20,024,384
Trainable params: 0
Non-trainable params: 20,024,384

None
Model: "sequential_19"

Layer (type)	Output Shape	Param #
vgg19 (Functional)	(None, 1, 1, 512)	20024384
flatten_7 (Flatten)	(None, 512)	0
dense_21 (Dense)	(None, 128)	65664
dense_22 (Dense)	(None, 64)	8256
dense_23 (Dense)	(None, 10)	650

Total params: 20,098,954
Trainable params: 74,570
Non-trainable params: 20,024,384


None

```
import pickle
import matplotlib.pyplot as plt
```

```
# Load the training history data from a pickle file
with open('VGG19classifier.pkl', 'rb') as f:
    trackLossClassifierVgg19 = pickle.load(f)
```

```
# Plot the validation and training loss curves over time
plt.plot(trackLossClassifierVgg19['val_loss'])
plt.plot(trackLossClassifierVgg19['loss'])
```

```
plt.legend(['Validation Loss', 'Train Loss'])
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.show()
```



```
modVgg19 = load_model("VGG19Classifier.h5")
val_vgg19 = (modVgg19.evaluate(x_test_final,y_test_labelled))[1]*100
print("Training accuracy is "+str((modVgg19.evaluate(x_train_final,y_train_labelled))[1]*100))
print("Test accuracy is "+str((modVgg19.evaluate(x_test_final,y_test_labelled))[1]*100))

313/313 [=====] - 11s 36ms/step - loss: 0.4882 - accuracy: 0.8282
1875/1875 [=====] - 60s 32ms/step - loss: 0.3922 - accuracy: 0.8569
Training accuracy is 85.69166660308838
313/313 [=====] - 10s 33ms/step - loss: 0.4882 - accuracy: 0.8282
Test accuracy is 82.81999826431274
```

```
gc.collect()
```

```
7954
```

▼ ResNet50V2

```
modResNet50V2 = Sequential()
modResNet50V2.add(ResNet50V2(include_top=False,weights='imagenet',input_shape=(56,56, 3)))
for i in range(len(modResNet50V2.layers)):
    modResNet50V2.layers[i].trainable = False
print(modResNet50V2.summary())
modResNet50V2.add(Flatten())
modResNet50V2.add(Dense(128, activation='relu'))
modResNet50V2.add(Dense(64, activation='relu'))
modResNet50V2.add(Dense(10, activation='softmax'))
print(modResNet50V2.summary())
modResNet50V2.compile(optimizer = 'adam',loss="categorical_crossentropy",metrics=[ 'accuracy'])
```

```
Model: "sequential_20"
```

Layer (type)	Output Shape	Param #
resnet50v2 (Functional)	(None, 2, 2, 2048)	23564800

```
=====
Total params: 23,564,800
Trainable params: 0
Non-trainable params: 23,564,800
```


```
None
Model: "sequential_20"
```

Layer (type)	Output Shape	Param #
resnet50v2 (Functional)	(None, 2, 2, 2048)	23564800
flatten_8 (Flatten)	(None, 8192)	0
dense_24 (Dense)	(None, 128)	1048704
dense_25 (Dense)	(None, 64)	8256
dense_26 (Dense)	(None, 10)	650

```
=====
Total params: 24,622,410
Trainable params: 1,057,610
Non-trainable params: 23,564,800
```


None

```
with open('ResNet50V2Classifier.pkl','rb') as data:
    trackLossClasifierResNet50V2 = pickle.load(data)
plt.plot(trackLossClasifierResNet50V2['val_loss'])
plt.plot(trackLossClasifierResNet50V2['loss'])
plt.legend(['Validation Loss', 'Train Loss'])
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.show()
modResNet50V2 = load_model("ResNet50V2Classifier.h5")
val_ResNet50V2 = (modResNet50V2.evaluate(x_test_final,y_test_labelled))[1]*100
print("Training accuracy is "+str((modResNet50V2.evaluate(x_train_final,y_train_labelled))[1]*100))
print("Test accuracy is "+str(val_ResNet50V2))
```



```
313/313 [=====] - 36s 115ms/step - loss: 0.6195 - accuracy: 0.7887
1875/1875 [=====] - 205s 109ms/step - loss: 0.5268 - accuracy: 0.8132
Training accuracy is 81.32166862487793
Test accuracy is 78.86999845504761
```

▼ MobileNet

```
modMobileNet = Sequential()
modMobileNet.add(MobileNet(include_top=False,weights='imagenet',input_shape=(56,56, 3)))
for i in range(len(modMobileNet.layers)):
    modMobileNet.layers[i].trainable = False
print(modMobileNet.summary())
modMobileNet.add(Flatten())
modMobileNet.add(Dense(128, activation='relu'))
modMobileNet.add(Dense(64, activation='relu'))
modMobileNet.add(Dense(10, activation='softmax'))
print(modMobileNet.summary())
modMobileNet.compile(optimizer = 'adam',loss="categorical_crossentropy",metrics=['accuracy'])
```

WARNING:tensorflow:`input_shape` is undefined or non-square, or `rows` is not in [128, 160, 192, 224]. Weights for input shape (224, 224, 3) will be automatically resized to match the input shape. This will affect class distances. Model: "sequential_21"

Layer (type)	Output Shape	Param #
mobilenet_1.00_224 (Functional)	(None, 1, 1, 1024)	3228864


Total params: 3,228,864
Trainable params: 0
Non-trainable params: 3,228,864

None
Model: "sequential_21"

Layer (type)	Output Shape	Param #
mobilenet_1.00_224 (Functional)	(None, 1, 1, 1024)	3228864
flatten_9 (Flatten)	(None, 1024)	0
dense_27 (Dense)	(None, 128)	131200
dense_28 (Dense)	(None, 64)	8256
dense_29 (Dense)	(None, 10)	650

Total params: 3,368,970
Trainable params: 140,106
Non-trainable params: 3,228,864

```
with open('MobileNetClassifier.pkl','rb') as data:
    trackLossClasifierMobileNet = pickle.load(data)
plt.plot(trackLossClasifierMobileNet['val_loss'])
plt.plot(trackLossClasifierMobileNet['loss'])
plt.legend(['Validation Loss', 'Train Loss'])
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.show()
modMobileNet = load_model("MobileNetClassifier.h5")
val_MobileNet = (modMobileNet.evaluate(preprocess_input(x_test_final),y_test_labelled))[1]*100
print("Training accuracy is "+str((modMobileNet.evaluate(preprocess_input(x_train_final),y_train_labelled))[1]*100))
print("Test accuracy is "+str(val_MobileNet))
```



```
313/313 [=====] - 20s 65ms/step - loss: 1.0337 - accuracy: 0.6310
1875/1875 [=====] - 120s 64ms/step - loss: 0.9794 - accuracy: 0.6532
Training accuracy is 65.3166651725769
Test accuracy is 63.099998235702515
```

EfficientNetB0

```
from tensorflow.keras.applications import EfficientNetB0
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Flatten, Dense

# Initialize a sequential model in Keras
modEfficientNetB0 = Sequential()

# Add the EfficientNetB0 model with pretrained weights from ImageNet, and freeze all of its layers
modEfficientNetB0.add(EfficientNetB0(include_top=False,weights='imagenet',input_shape=(56,56, 3)))
for i in range(len(modEfficientNetB0.layers)):
    modEfficientNetB0.layers[i].trainable = False

# Print a summary of the model architecture
print(modEfficientNetB0.summary())

# Add a flattening layer to convert the output of the EfficientNetB0 model into a 1D array
modEfficientNetB0.add(Flatten())

# Add a fully connected layer with 128 neurons and ReLU activation
modEfficientNetB0.add(Dense(128, activation='relu'))

# Add a second fully connected layer with 64 neurons and ReLU activation
modEfficientNetB0.add(Dense(64, activation='relu'))

# Add a final fully connected layer with 10 neurons and softmax activation
modEfficientNetB0.add(Dense(10, activation='softmax'))

# Print a summary of the model architecture
print(modEfficientNetB0.summary())

# Compile the model using the Adam optimizer, categorical crossentropy loss, and accuracy as a metric
modEfficientNetB0.compile(optimizer = 'adam',loss="categorical_crossentropy",metrics=['accuracy'])
```

Model: "sequential_22"


Layer (type)	Output Shape	Param #
efficientnetb0 (Functional)	(None, 2, 2, 1280)	4049571
Total params: 4,049,571		

Trainable params: 0
Non-trainable params: 4,049,571

None		
Model: "sequential_22"		
Layer (type)	Output Shape	Param #
=====		
efficientnetb0 (Functional)	(None, 2, 2, 1280)	4049571
flatten_10 (Flatten)	(None, 5120)	0
dense_30 (Dense)	(None, 128)	655488
dense_31 (Dense)	(None, 64)	8256
dense_32 (Dense)	(None, 10)	650
=====		
Total params: 4,713,965		
Trainable params: 664,394		
Non-trainable params: 4,049,571		
None		

```
with open('EfficientNetB0Classifier.pkl','rb') as data:
    trackLossClassifierEfficientNetB0 = pickle.load(data)

plt.plot(trackLossClasifierEfficientNetB0['val_loss'])
plt.plot(trackLossClasifierEfficientNetB0['loss'])
plt.legend(['Validation Loss', 'Train Loss'])
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.show()
modEfficientNetB0.load_weights("EfficientNetB0Classifier.h5")
val_EfficientNetB0 = (modEfficientNetB0.evaluate(x_test_final,y_test_labelled))[1]*100
print("Training accuracy is "+str((modEfficientNetB0.evaluate(x_train_final,y_train_labelled))[1]*100))
print("Test accuracy is "+str(val_EfficientNetB0))
```



```
313/313 [=====] - 51s 162ms/step - loss: 0.3369 - accuracy: 0
1875/1875 [=====] - 301s 161ms/step - loss: 0.2578 - accuracy
Training accuracy is 90.41333198547363
Test accuracy is 87.80999779701233
```

```
print("The accuracy for VGG16 "+str(val_vgg16))
print("The accuracy for VGG19 "+str(val_vgg19))
print("The accuracy for ResNet50V2 "+str(val_ResNet50V2))
print("The accuracy for MobileNet "+str(val_MobileNet))
print("The accuracy for EfficientNetB0 "+str(val_EfficientNetB0))
```

```
The accuracy for VGG16 82.81000256538391
The accuracy for VGG19 82.81999826431274
The accuracy for ResNet50V2 78.86999845504761
The accuracy for MobileNet 63.099998235702515
The accuracy for EfficientNetB0 87.80999779701233
```

+ Code + Text